Solution Brief

Cross-Industry Architecture
Intel® Cloud Optimization Modules

intel software

Intel® Cloud Optimization Modules for GCP*: nanoGPT Distributed Training

Author: Benjamin Consolvo

Date: December 6, 2023

Fine-tune nanoGPT in a distributed architecture on GCP with Intel® Xeon® Scalable Processors

The **Intel® Cloud Optimization Modules for GCP*: nanoGPT Distributed Training** is designed to illustrate the process of fine-tuning a large language model (LLM) with 4th Generation **Intel® Xeon® Scalable Processors** on **Google Cloud Platform (GCP)***. Specifically, we show the process of fine-tuning a nanoGPT model with 124M parameters on the **OpenWebText dataset** in a distributed architecture. The project builds upon the initial codebase of **nanoGPT** built by Andrej Karpathy. The objective is to understand how to set up a distributed system so that you can fine-tune the model to your specific workload. The result of this module will be a base LLM that can generate words, or tokens, that will be suitable for your use case when you modify it to your specific objective and dataset.

Use it as a reference solution for:

- Setting up an GCP cluster for distributed training.
- Fine-tuning an LLM on a single machine.
- Fine-tuning an LLM in a distributed system, taking advantage of Intel optimizations.

Who needs it?

- Developers aiming to fine-tune their LLMs on multiple Intel Xeon CPUs, leveraging Intel's accelerated deep learning software libraries, including **Intel® Extension for PyTorch*** and **Intel® oneAPI Collective Communications Library (oneCCL)**.
- Developers interested in learning the process of setting up clusters for distributed training.

What it does

This module demonstrates how to transform a standard single-node PyTorch training scenario into a high-performance distributed training scenario across multiple CPUs. To fully capitalize on Intel hardware and further optimize the fine-tuning process, this module integrates the **Intel® Extension for PyTorch*** and **Intel® oneAPI Collective Communications Library (oneCCL)**. The module serves as a guide to setting up a cluster for distributed training while showcasing a complete project for fine-tuning LLMs.

- It provides step-by-step instructions for configuring a cluster, simplifying the process of establishing a distributed training environment.
- It serves as a guide through the entire lifecycle of fine-tuning LLMs, starting from data preprocessing to model fine-tuning.
- The module capitalizes on Intel® Extension for PyTorch, harnessing the power of **Intel® Advanced Matrix Extension (Intel® AMX)** instruction sets. This enables significant acceleration of the fine-tuning process, boosting overall training performance. The use of Intel's optimized communications library, oneCCL, ensures that distributed workflows are streamlined, enhancing efficiency in a multi-node training setup.

In summary, this module empowers you to harness the full potential of Intel hardware for distributed training an LLM.

Cloud Solution Architecture

To form the cluster, the cloud solution implements GCP virtual machines from the **C3** series. To enable seamless communication between the instances, each of the machines are connected to the same virtual network and a permissive network security group is established that allows all traffic from other nodes within the cluster. The raw dataset is taken from Hugging Face*, and once the model has been trained, the weights are saved to the virtual machines (Figure 1).

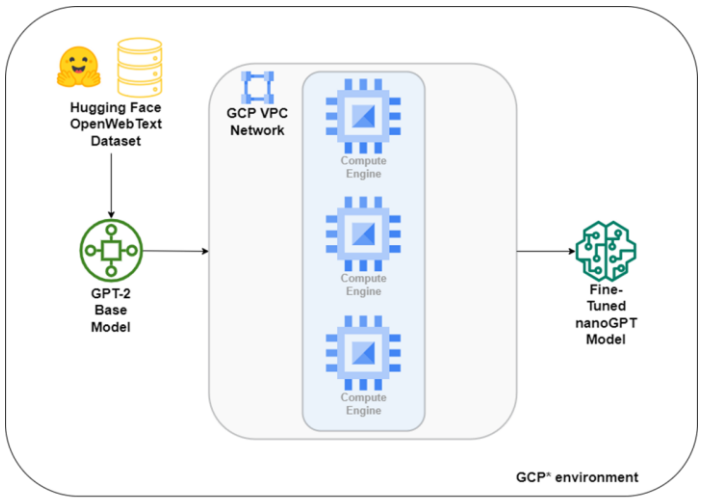

Figure 1: Architectural Diagram of GCP nanoGPT Distributed Training module. Image by author.

Code Highlights

Enable the Intel Extension for PyTorch

The Intel Extension for PyTorch elevates PyTorch performance on Intel hardware with the integration of the newest features and optimizations that have not yet been incorporated into open source PyTorch. This extension efficiently utilizes Intel hardware capabilities including Intel AMX. Unleashing this power is straightforward – just wrap your model and optimizer objects with `ipex.optimize`.

```
# Set up CPU autocast and bfloat16 dtype
dtype = torch.bfloat16
self.autocast_ctx_manager = torch.cpu.amp.autocast(
        cache_enabled=True, dtype=dtype
)

# Wrap both Pytorch model and Optimizer
self.model, self.optimizer = ipex.optimize(
        self.model, optimizer=self.optimizer,
        dtype=dtype, inplace=True, level="O1",
)
```

Commented [EJ1]: Perhaps we should be clearer w
something like "the newest features and optimizati
that have not yet been incorporated into open sou
PyTorch"

Gradient Accumulation with Hugging Face Accelerate

The Accelerate library by Hugging Face streamlines the gradient accumulation process. This package helps to abstract away the complexity of supporting multi-CPUs/GPUs and provides an intuitive, user-friendly API, making gradient accumulation and clipping hassle-free during the training process.

```
# Initializing Accelerator object
self.accelerator = Accelerator(
        gradient_accumulation_steps=gradient_accumulation_steps,
        cpu=True,
)

# Gradient Accumulation
with self.accelerator.accumulate(self.model):
       with self.autocast_ctx_manager:
              _, loss = self.model(X, Y)
       self.accelerator.backward(loss)
       loss = loss.detach() / gradient_accumulation_steps

# Gradient Clipping
self.accelerator.clip_grad_norm_(
       self.model.parameters(), self.trainer_config.grad_clip
)
```

Distributed Training

For distributed training, we utilized oneCCL. With optimized communication patterns, oneCCL enables developers and researchers to train newer and deeper models more quickly across multiple nodes. It offers a tool called `mpirun`, which allows you to seamlessly launch distributed training workloads.

```
# Generating Multi-CPU config
accelerate config --config_file ./multi_config.yaml

# Launching Distributed Training job
mpirun -f ~/hosts -n 3 -ppn 1 -genv LD_PRELOAD="/usr/lib/x86_64-linux-gnu/libtcmalloc.so" accelerate
launch --config_file ./multi_config.yaml main.py
```

Next Steps

Download the module from GitHub ›

Check out the full suite of Intel® Cloud Optimization Modules ›

Register for office hours for implementation support from Intel engineers ›

Come chat with us on our DevHub Discord server to keep interacting with other developers ›